

# Deep learning Feed-forward neural networks

Ole Winther

Bioinformatics Centre, Department of Biology  
University of Copenhagen (UCph)

Dept for Applied Mathematics and Computer Science  
Technical University of Denmark (DTU)



November 18, 2025

# Objectives of lecture

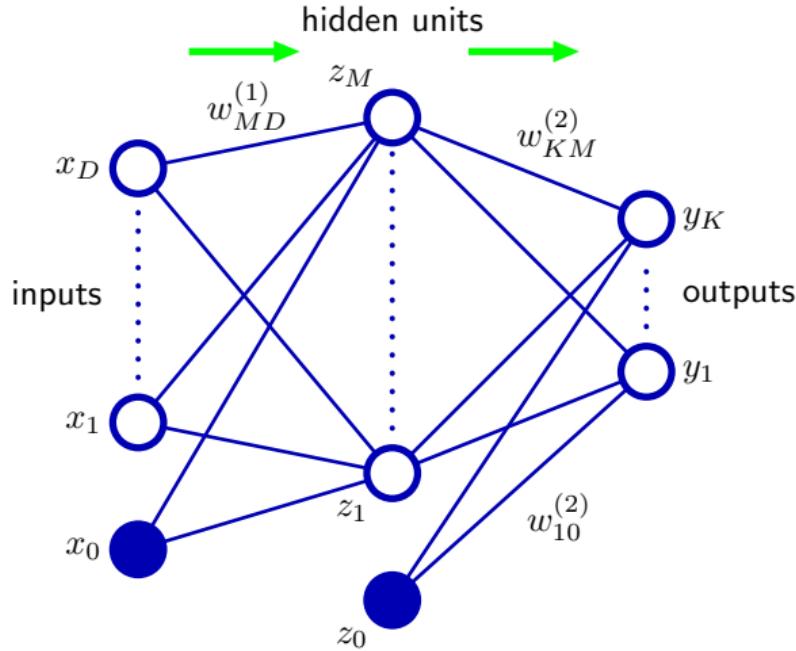
- Feed-forward neural network (FFNN)
- Next week:
- Training with **error back-propagation**
- We only need to understand the principles
- **Autograd** - automated differentiation handles the derivation for us!



Many thanks to Tapani Raiko for making and sharing first version of these slides!

# Part 1: Feed-forward neural networks

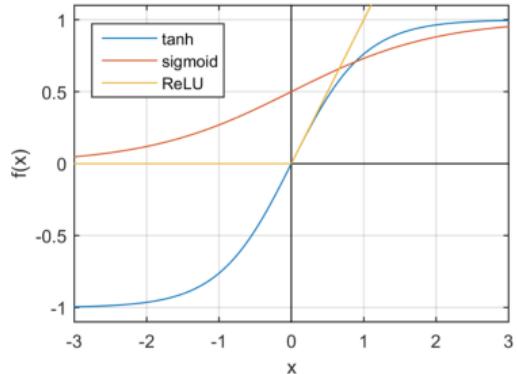
# Feed forward neural networks



$$y_k(\mathbf{x}, \mathbf{w}) = \sigma \left( \sum_{j=0}^M w_{kj}^{(2)} f \left( \underbrace{\sum_{i=0}^D w_{ji}^{(1)} x_i}_{z_j} \right) \right)$$

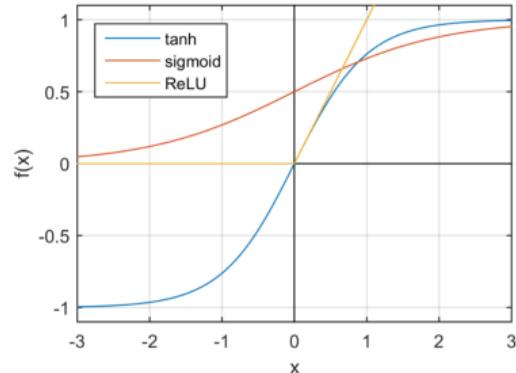
# Non-linearity and training

- Linear activation functions will give a linear network.
- Logistic function  $\sigma(a) = \frac{1}{1+e^{-a}}$
- Hyperbolic tangent  $\tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$
- Rectified linear  $\text{relu}(a) = \max(0, a)$



# Non-linearity and training

- Linear activation functions will give a linear network.
- Logistic function  $\sigma(a) = \frac{1}{1+e^{-a}}$
- Hyperbolic tangent  $\tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$
- Rectified linear  $\text{relu}(a) = \max(0, a)$



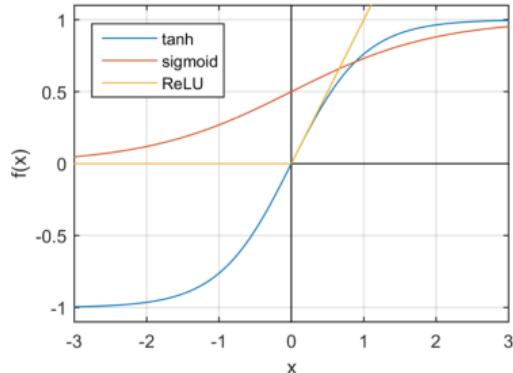
- Supervised learning
- Labeled training set

$$\mathcal{D} = \{(x_i, t_i) | i = 1, \dots, n\} .$$

- Input  $x_i$  and target  $t_i$ .

# Non-linearity and training

- Linear activation functions will give a linear network.
- Logistic function  $\sigma(a) = \frac{1}{1+e^{-a}}$
- Hyperbolic tangent  $\tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$
- Rectified linear  $\text{relu}(a) = \max(0, a)$



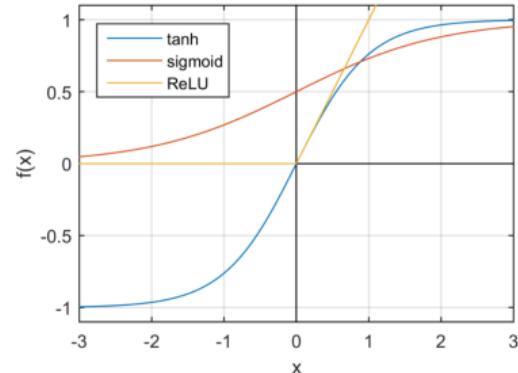
- Supervised learning
- Labeled training set

$$\mathcal{D} = \{(x_i, t_i) | i = 1, \dots, n\} .$$

- Input  $x_i$  and target  $t_i$ .
- Minimize training error by (stochastic) gradient descent

# Non-linearity and training

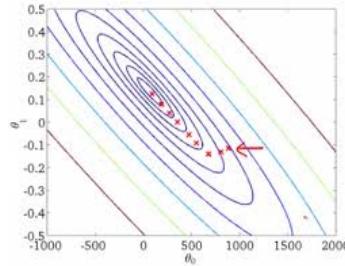
- Linear activation functions will give a linear network.
- Logistic function  $\sigma(a) = \frac{1}{1+e^{-a}}$
- Hyperbolic tangent  $\tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$
- Rectified linear  $\text{relu}(a) = \max(0, a)$



- Supervised learning
- Labeled training set

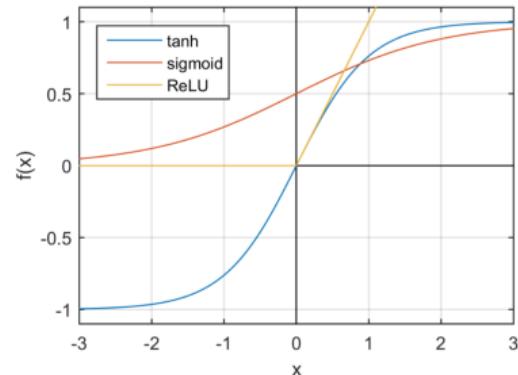
$$\mathcal{D} = \{(x_i, t_i) | i = 1, \dots, n\} .$$

- Input  $x_i$  and target  $t_i$ .
- Minimize training error by (stochastic) gradient descent



# Non-linearity and training

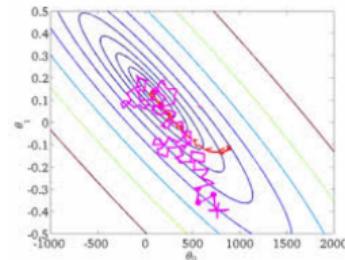
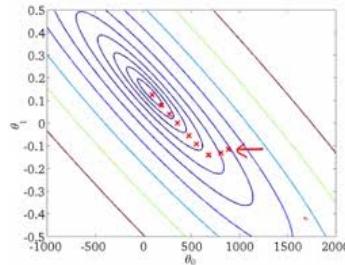
- Linear activation functions will give a linear network.
- Logistic function  $\sigma(a) = \frac{1}{1+e^{-a}}$
- Hyperbolic tangent  $\tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$
- Rectified linear  $\text{relu}(a) = \max(0, a)$



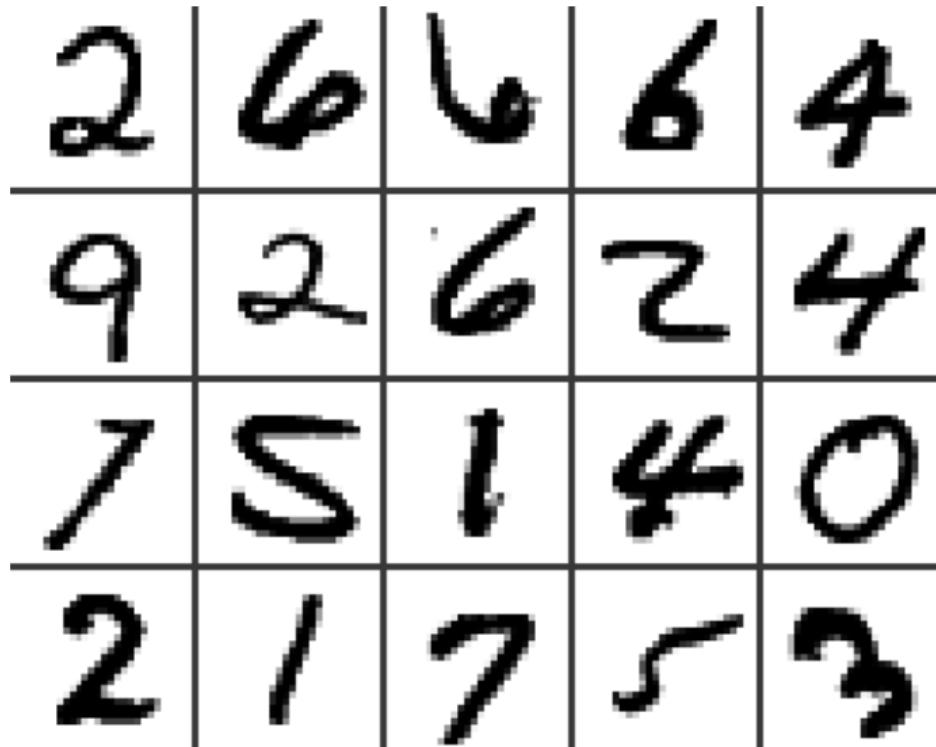
- Supervised learning
- Labeled training set

$$\mathcal{D} = \{(x_i, t_i) | i = 1, \dots, n\} .$$

- Input  $x_i$  and target  $t_i$ .
- Minimize training error by (stochastic) gradient descent



## Example: MNIST handwritten digits

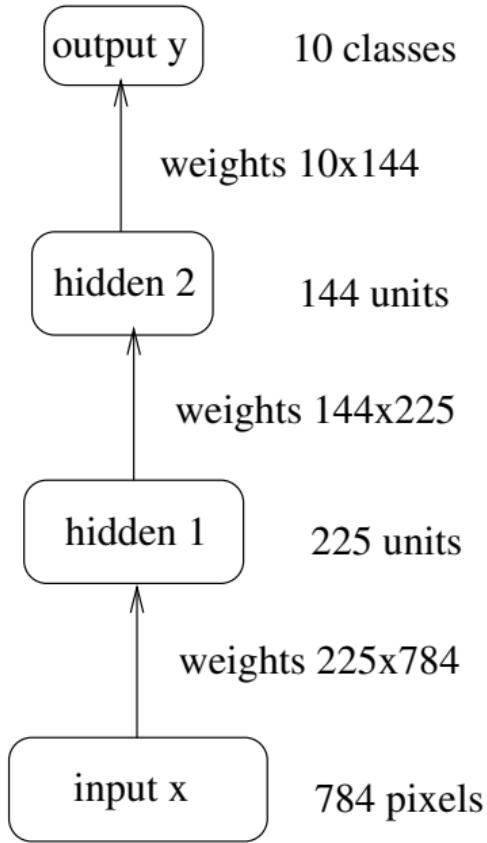


Train a network to classify  $28 \times 28$  images.

Data: 60000 input images  $\mathbf{x}(n)$  and labels  $t(n)$ .

Example model gives around 1.2% test error.

# Example Network



$$y = h^{(3)} = \text{softmax}(W^{(3)}h^{(2)} + b^{(3)})$$

$$h^{(2)} = \text{relu}(W^{(2)}h^{(1)} + b^{(2)})$$

$$h^{(1)} = \text{relu}(W^{(1)}x + b^{(1)})$$

$$\text{softmax}(\mathbf{z})_i = \frac{\exp(z_i)}{\sum_j \exp(z_j)}$$
$$\text{relu}(z) = \max(0, z)$$

# Softmax

- Softmax function

$$\text{softmax}(\mathbf{z})_i = \frac{\exp(z_i)}{\sum_j \exp(z_j)}$$

has two nice properties:

- $\text{softmax}(\mathbf{z})_i \geq 0$
- $\sum_i \text{softmax}(\mathbf{z})_i = 1$

# Softmax

- Softmax function

$$\text{softmax}(\mathbf{z})_i = \frac{\exp(z_i)}{\sum_j \exp(z_j)}$$

has two nice properties:

- $\text{softmax}(\mathbf{z})_i \geq 0$
- $\sum_i \text{softmax}(\mathbf{z})_i = 1$
- MNIST, output labels: 0, 1, ..., 9.
- Output of network

$$\mathbf{y} = \text{softmax}(\mathbf{W}^{(3)}\mathbf{h}^{(2)} + \mathbf{b}^{(3)})$$

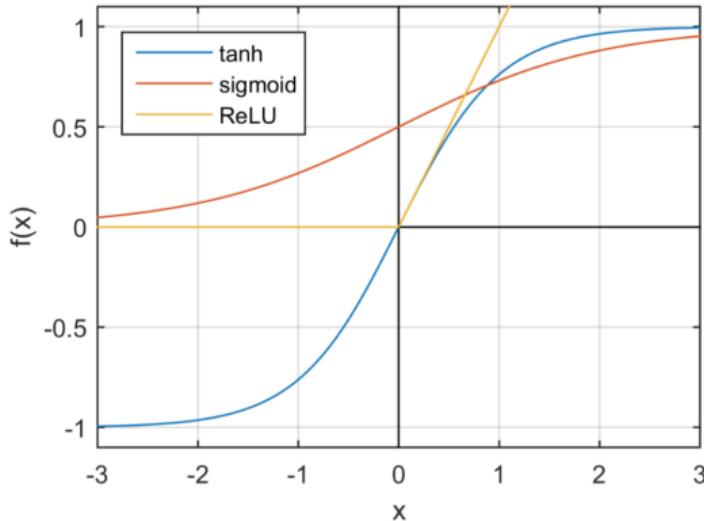
interpreted as class(-conditional) probabilities:

- So given input  $\mathbf{x}$ , according to the model, the probability of digit  $i$  is:

$$p(\text{digit} = i | \mathbf{x}) = y_{i+1}$$

- with  $i = 0, \dots, 9$ .

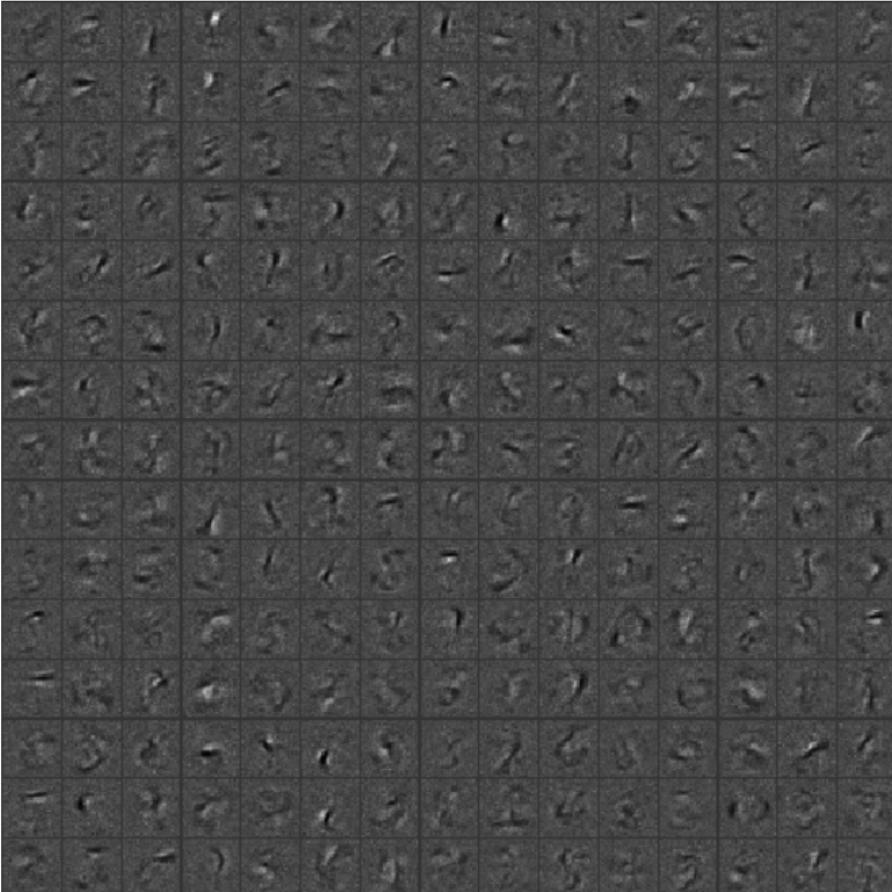
# On activation functions



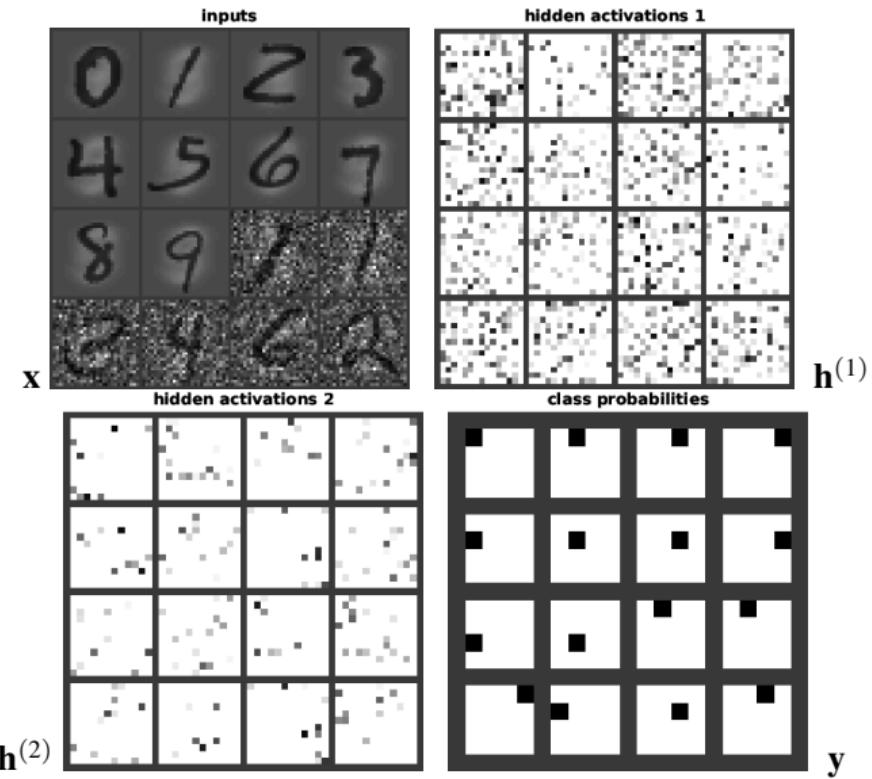
- $\text{relu}(z) = \max(0, z)$  is replacing old sigmoid and tanh.
- Note that identity function would lead into:

$$\begin{aligned}\mathbf{h}^{(2)} &= \mathbf{W}^{(2)}\mathbf{h}^{(1)} + \mathbf{b}^{(2)} \\ &= \mathbf{W}^{(2)}(\mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)}) + \mathbf{b}^{(2)} \\ &= (\mathbf{W}^{(2)}\mathbf{W}^{(1)})\mathbf{x} + (\mathbf{W}^{(2)}\mathbf{b}^{(1)} + \mathbf{b}^{(2)}) \\ &= \mathbf{W}'\mathbf{x} + \mathbf{b}'\end{aligned}$$

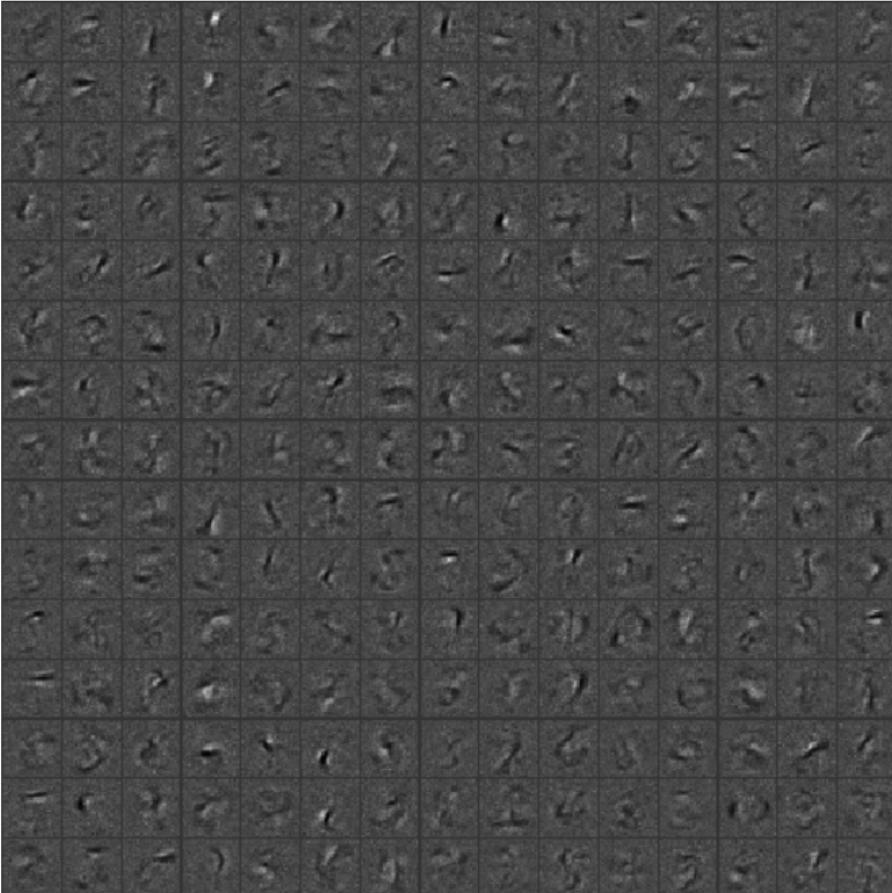
Weight matrix  $\mathbf{W}^{(1)}$  size  $225 \times 784$



Signals  $\mathbf{x} \rightarrow \mathbf{h}^{(1)} \rightarrow \mathbf{h}^{(2)} \rightarrow \mathbf{h}^{(3)}$



Weight matrix  $\mathbf{W}^{(1)}$  size  $225 \times 784$



## Part 2:

Doing yourself - understand how  
feed-forward neural networks  
approximate functions

# How a feed-forward neural network learns XOR

- Consider two-layer network with two hidden units:

$$y(\mathbf{x}) = \Theta \left( W_1^{(2)} h_1^{(1)} + W_2^{(2)} h_2^{(1)} + b_1^{(1)} \right)$$

$$h_1^{(1)} = \Theta \left( W_{11}^{(1)} x_1 + W_{12}^{(1)} x_2 + b_1^{(1)} \right)$$

$$h_2^{(1)} = \Theta \left( W_{21}^{(1)} x_1 + W_{22}^{(1)} x_2 + b_2^{(1)} \right)$$

- Step function activation function:  $\Theta(z) = 1$  if  $z \geq 0$  and 0 otherwise

# How a feed-forward neural network learns XOR

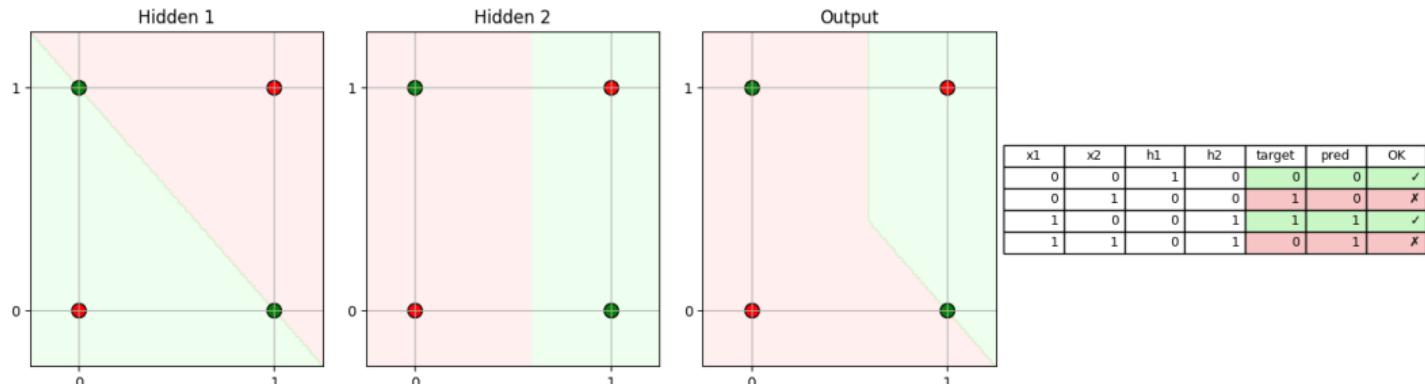
- Consider two-layer network with two hidden units:

$$y(\mathbf{x}) = \Theta\left(W_1^{(2)}h_1^{(1)} + W_2^{(2)}h_2^{(1)} + b_1^{(1)}\right)$$

$$h_1^{(1)} = \Theta\left(W_{11}^{(1)}x_1 + W_{12}^{(1)}x_2 + b_1^{(1)}\right)$$

$$h_2^{(1)} = \Theta\left(W_{21}^{(1)}x_1 + W_{22}^{(1)}x_2 + b_2^{(1)}\right)$$

- Step function activation function:  $\Theta(z) = 1$  if  $z \geq 0$  and 0 otherwise
- Fetch [XOR\\_NNSandbox.ipynb](#) from Absalon and turn the knobs to learn XOR



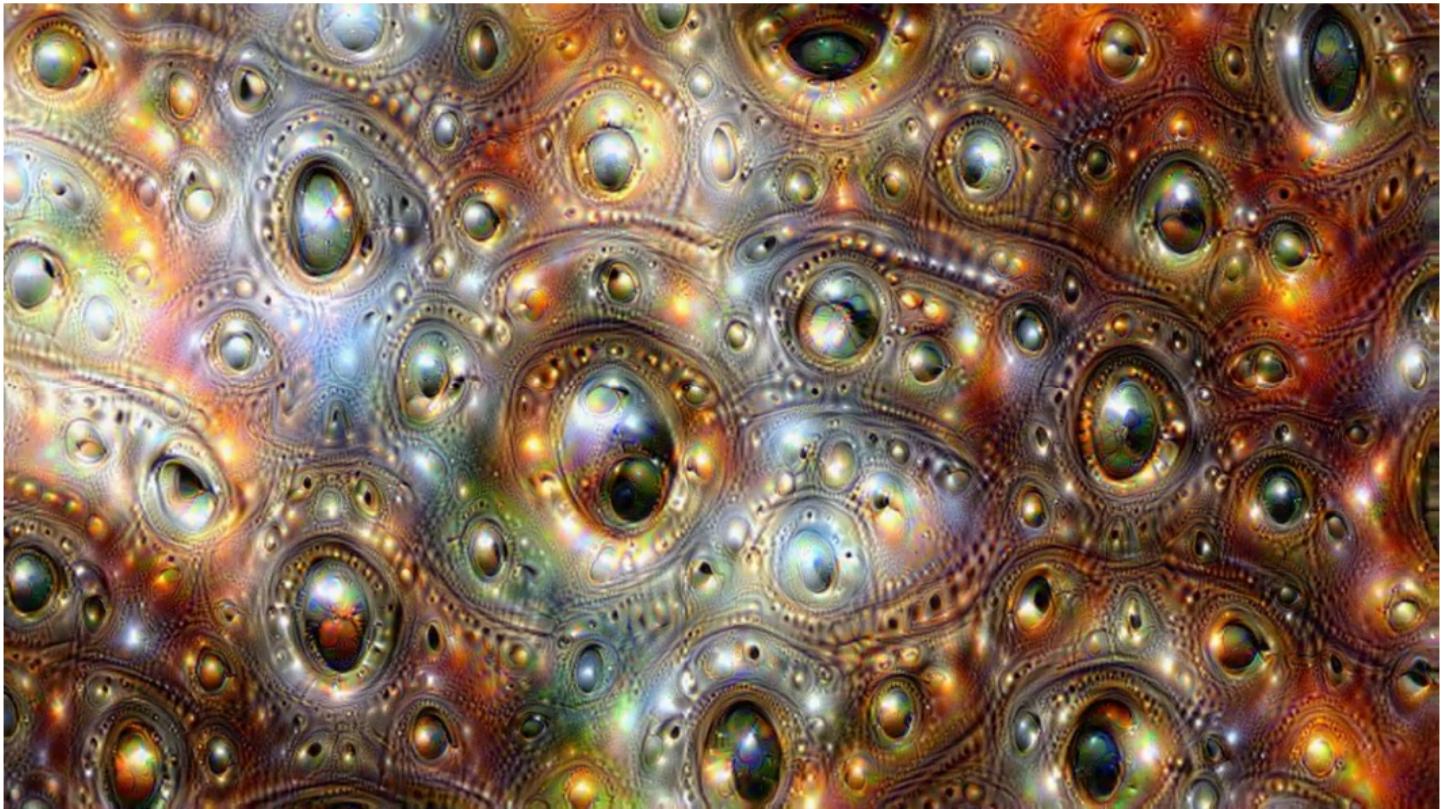
- Red is output 1 and green is output 0.
- We will also consider this problem in Assignment 1.

# TensorFlow Playground

[playground.tensorflow.org](https://playground.tensorflow.org)

# References

- Online book: [Michael Nielsen, Neural networks and deep learning](#)
- Book also online: [Ian Goodfellow and Yoshua Bengio and Aaron Courville, Deep learning](#)
- LeCun, Yann, Yoshua Bengio, and Geoffrey Hinton, Deep learning, *Nature* 521.7553 (2015): 436-444.
- Mnih, Volodymyr, et al., Human-level control through deep reinforcement learning, *Nature* 518.7540 (2015): 529-533.
- Alipanahi, Babak, et al., Predicting the sequence specificities of DNA-and RNA-binding proteins by deep learning, *Nature biotechnology* (2015).
- Silver, David, et al., Mastering the game of Go with deep neural networks and tree search, *Nature* 529.7587 (2016): 484-489.
- Xu, K., Ba, J., Kiros, R., Cho, K., Courville, A., Salakhutdinov, R., Zemel, R., & Bengio, Y. (2015), Show, attend and tell: Neural image caption generation with visual attention. arXiv preprint arXiv:1502.03044.
- Mansimov, Elman, et al., Generating Images from Captions with Attention. arXiv preprint arXiv:1511.02793 (2015).
- Larsen, Anders Boesen Lindbo, Soren Kaae Sonderby, and Ole Winther. Autoencoding beyond pixels using a learned similarity metric. arXiv preprint arXiv:1512.09300 (2015).
- Ian Goodfellow, Yoshua Bengio, and Aaron Courville, [Deep learning](#), 2016.
- Michael Nielsen, [Neural Networks and Deep Learning](#)
- Alex Krizhevsky, Ilya Sutskever, and Geoffrey E. Hinton, ImageNet Classification with Deep Convolutional, NIPS, 2012. Neural Networks,



Thanks!  
Ole Winther